

Justus Purat & Alexander Kammeyer
Software Project Distributed Systems

Consumption Data Forecast for HPC Systems

Wrap-up

M. Ch

A. Er

A. Huth

M. Karn

Y. Kaya

M. Zent

Overview

- Introduction
- Summary of the Sprints
- Presentation of the employed model
- Application overview and experimental observations
- Comment on Lag Llama
- Outlook

Project background

- Goal: Optimization of HPC clusters by making them more adaptive
 - rescheduling tasks to times where energy prices are low
 - choosing alternative energy sources
- Some datasources have a limited forecast
 - forecast energy prices and co2 emission using a ML model
- Train ML using time-series data from InfluxDB

Summary of the Sprints

Sprint 1

- Database connection
- Queried the database to get familiar with the structure and shape
- Analysing the data with python and pandas library
- Creating Charts
- Analysed the results from the Group of the last year

Summary of the Sprints

Sprint 2

- Looked for correlations
- Feature Engineering
- Made thoughts about the Model we will use
- worked on
 - Random Forest Regression
 - Prophet
 - Lag Lama

Summary of the Sprints

Sprint 3

- Short Sprint
- Worked further on the models, esp. LSTM and Lag Llama

Summary of the Sprints

Sprint 4

- Finished the preprocessing part
- Included Max Pooling Layers in LSTM
- Were able to predict 2h into the future
- Started with the Report

Data Handling

Feature Engineering:

- Finalized the Univariate approach.
 - because of the uncertainty of availability of correlated features at the time of forecasting.
 - all other state-of-the-art forecasting models use similar univariate approach.
- Found lagged values of price in the past which correlated highly with current price values.
 - best lagged features were within the 1 week period in the past.
- Created rolling means over the windows of upto 1 week from the past; same as lagged features.
- Also used sine and cosine transformations for time based features to cater for their cyclic nature.

Preprocessing:

- After many experimentations, found that 48 hours sequence in the past to predict 1 hour in the future give the best results.
- Training and testing splits (90% training & 10% testing).
- Applied MinMax scalers on features and labels to normalize their values.

Loss Function

Drawbacks of mse:

- Quadratic penalties for large errors.
- Makes the model more conservative by finding a middle ground that minimizes squared error across all data points.
- Symmetry; all errors are treated the same; no domain knowledge can be embedded.

Custom Loss function:

- Value of delta controls the behavior of the loss function.
- Quadratic for errors less than delta and linear for errors greater than delta.
- Asymmetric nature - we can customize it according to our domain needs e.g. It is better to underpredict than over predict. Over predictions incur 30% more penalty than under predictions.
- It also gives us more control through adjustment of penalties for peaks and dips.
- Penalty increases for under prediction as y_{true} (normalized price) gets higher for peaks. Similarly, penalty increase for over prediction if y_{true} gets lower for downward spikes.

Model Architecture

1D Convolution and Max pooling:

- 1D convolution applies filters over 1 dimensional temporal sequences.
- The filters learn local patterns and features across different time steps within a sequence.
- Filters are sliding windows that detect specific "shapes" in data (e.g., a sudden spike, a dip, a short-term trend reversal).
- Max pooling reduces the size of sequences while retains all the important information.

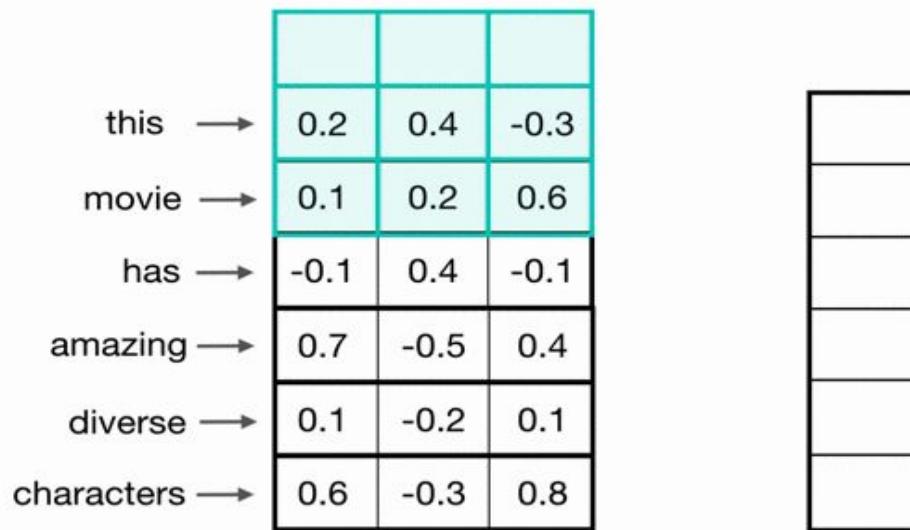
LSTM:

- The CNN extracts robust, high-level local features, and then the LSTM processes these learned features sequentially to understand long-term dependencies.
- The combinations allows the LSTM to receive a richer, more abstract representation of the sequence, rather than just raw time series values.

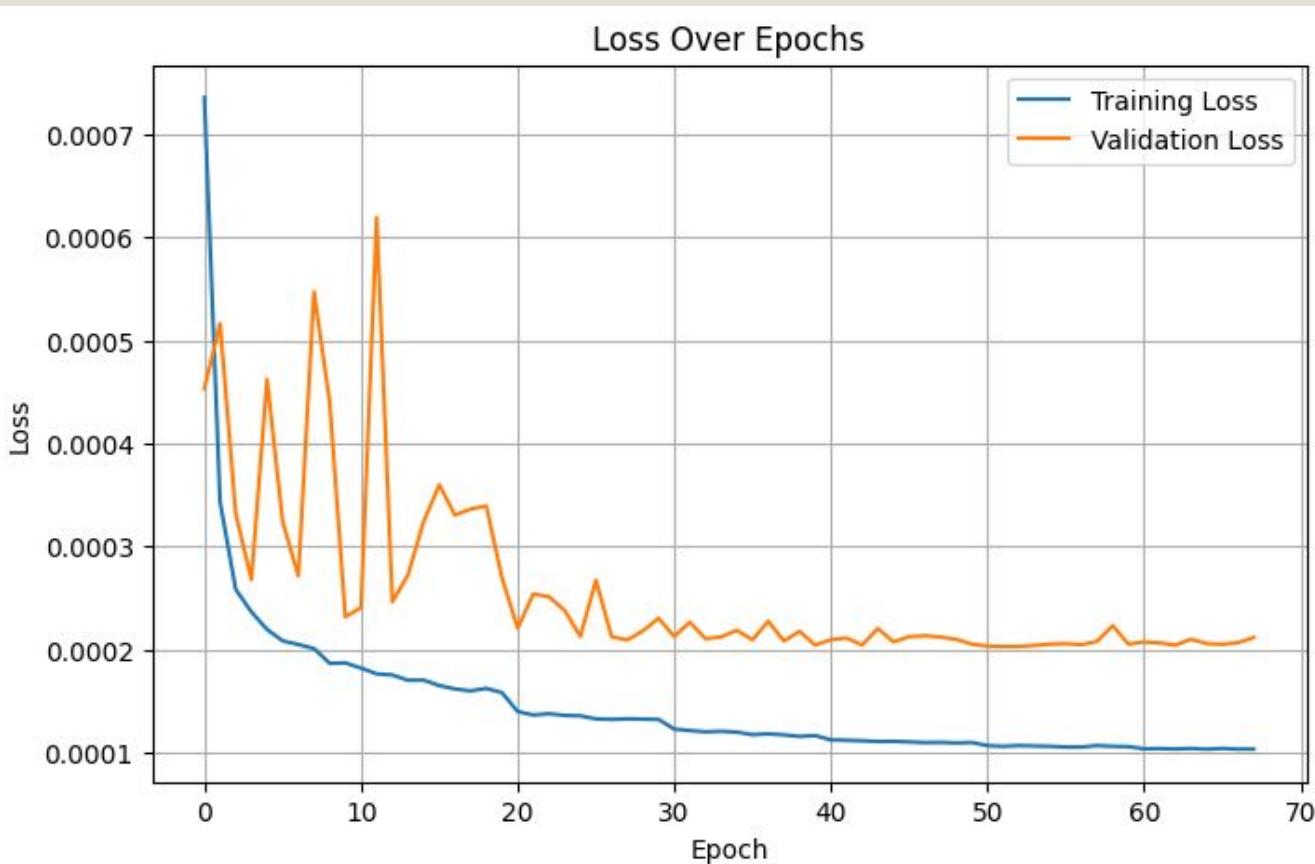
ANN:

- The tail of the model is a simple neural networks which gives us our forecast for 1 hour in the future.

Abstract Representation of 1D Convolution



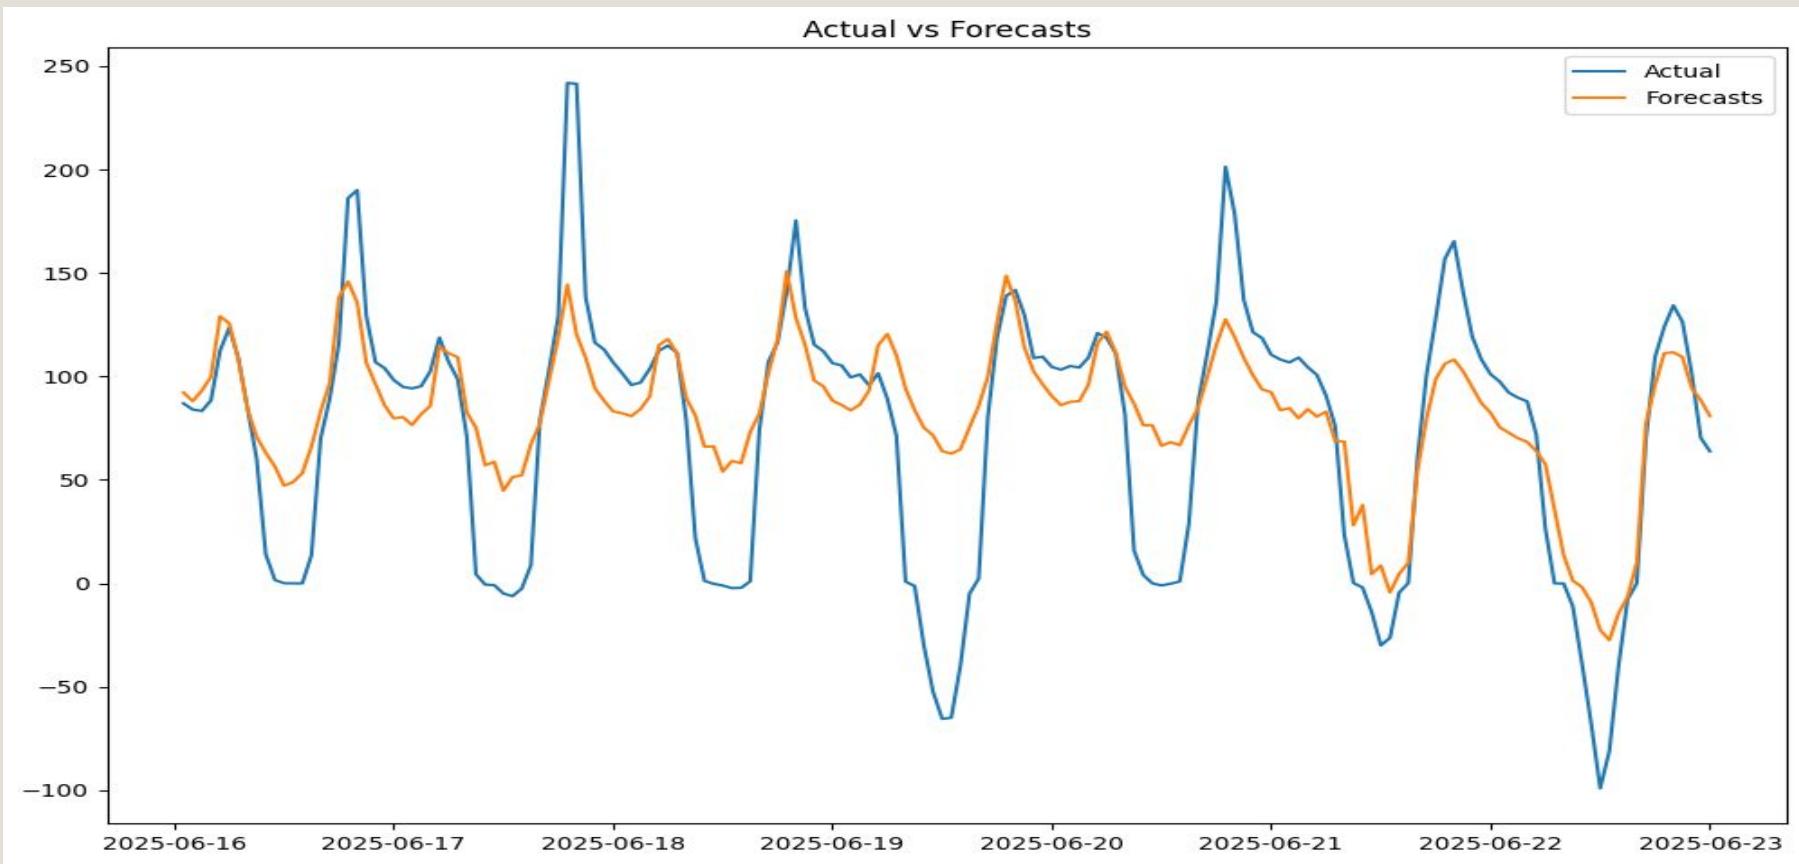
Training with Custom Loss



Achieved best results
at Epoch 53

loss	1.0905e-04
mae	0.0060
mse	1.0004e-04
val_loss	2.0292e-04
val_mae	0.0087
val_mse	3.1625e-04

Rolling forecasts (Recursively)



Rolling forecasts (Continued)

Errors representation:

- Mean Squared Error: 1767.57
- Mean Absolute Error: 31.02

What this tells us is on average a forecasted price value is 31 units away from the actual value.

This however, is the error representation if we forecast for 1 week in the future. This error significantly declines for shorter forecast horizons as Michael will discuss later.

Application Structure

ConsumptionData.py

- Data fetching and handling
- Distance correlation analysis
- Transformation to and from gradient values
- Outlier removal
- Adding lagged and rolling mean features

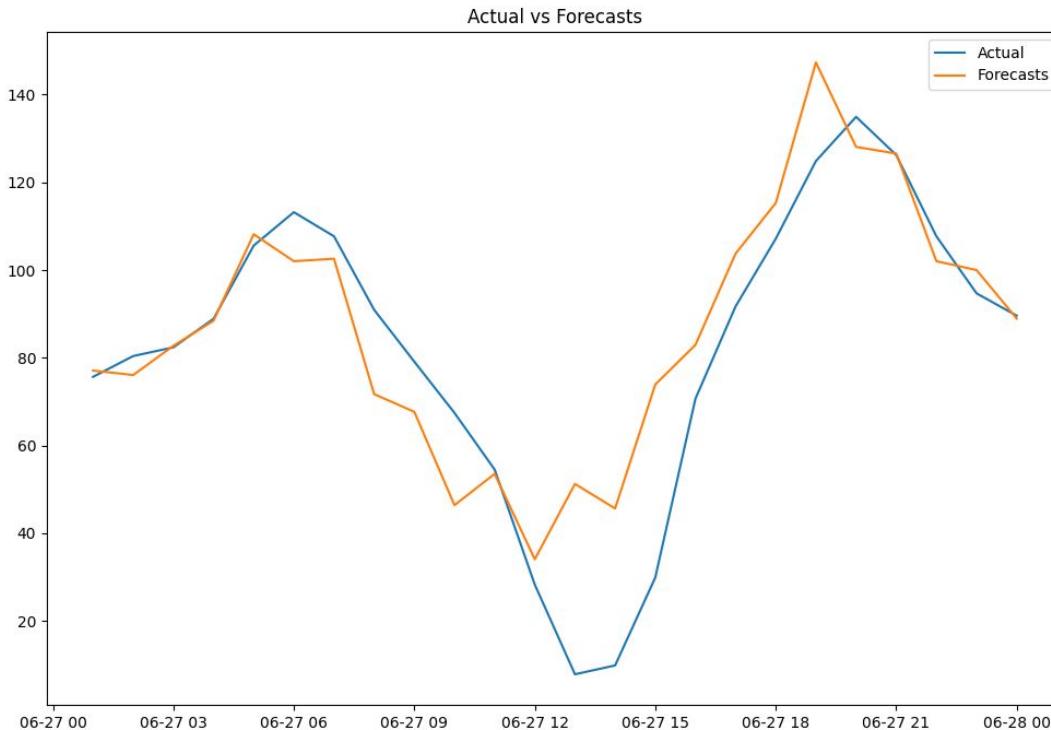
modelLSTM.py & TwoSidedAsymmetricHuberLoss.py

- Implement the aforementioned LSTM model, and the employed loss function

training.py & forecast.py

- Call the training and forecasting routines after corresponding pre-processing

Experiments on Forecasting a whole Day



Training data

3 years,
ending 31.12.2024

Absolute values

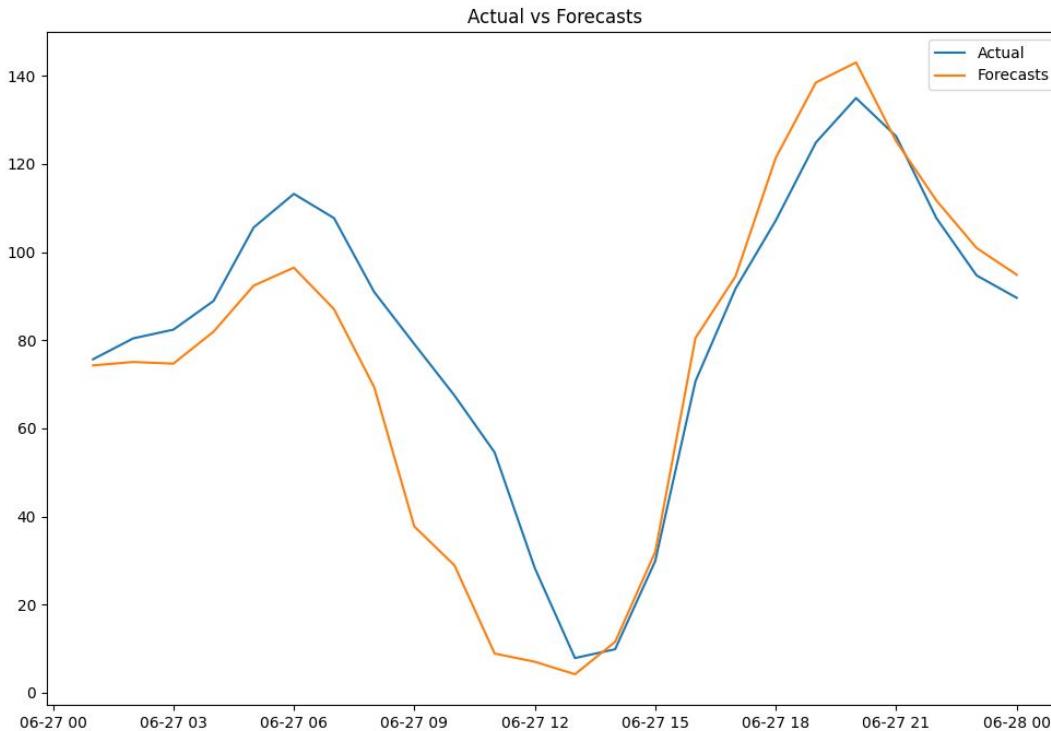
No outlier removal

Artificial features

47 lagged and
rolling-mean features
by $dCor > 0.8$

Normalized MAE
9.2%

Experiments on Forecasting a whole Day



Training data

3 years,
ending 31.05.2025

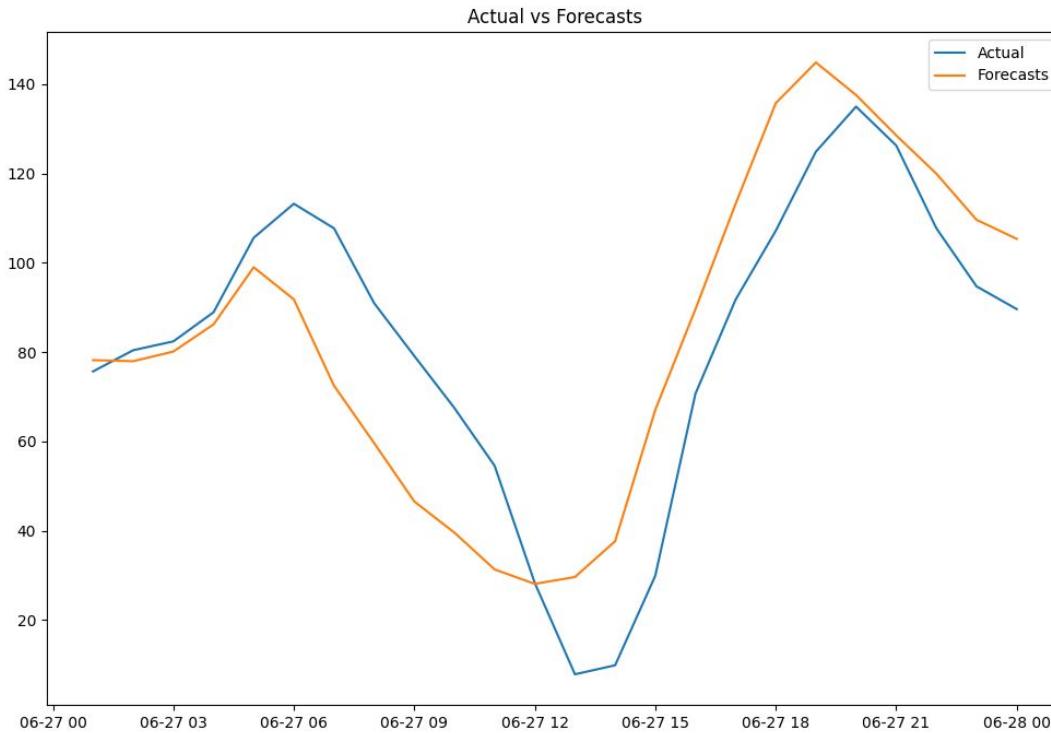
Absolute values

No outlier removal

Artificial features
47 lagged and
rolling-mean features
by $dCor > 0.8$

Normalized MAE
10.3%

Experiments on Forecasting a whole Day



Training data

3 years,
ending 31.12.2024

Gradient values

Outlier removal
z-score at 3 s

Artificial features

39 lagged and
rolling-mean features
by dCor > 0.8

Normalized MAE
11.9%

Observations

- 24h-forecasts achieve a normalized MAE of 9-13% at yearly or monthly retraining
- The closer the training data's end-date to the date for which the forecast has to be done, the worse the normalized MAE is in total, but can become better on certain time sections
- Offset between forecast and actual data of about an hour
- Using gradient data with outlier removal
 - Tends to be more stable, i.e. to preserve the actual curve's form
 - Reduces the number of needed artificial features
 - But performs slightly worse in normalized MAE

Forecasting with Foundation Model

Goal:

- Utilize foundation models in forecasting energy price mix and carbon dioxide emission.

Fundamental Idea: A time series foundation model can create forecasts without pretraining, similar to how a large language model (LLM) can output task without being pre-trained on a task.

Tool: Lag Llama¹

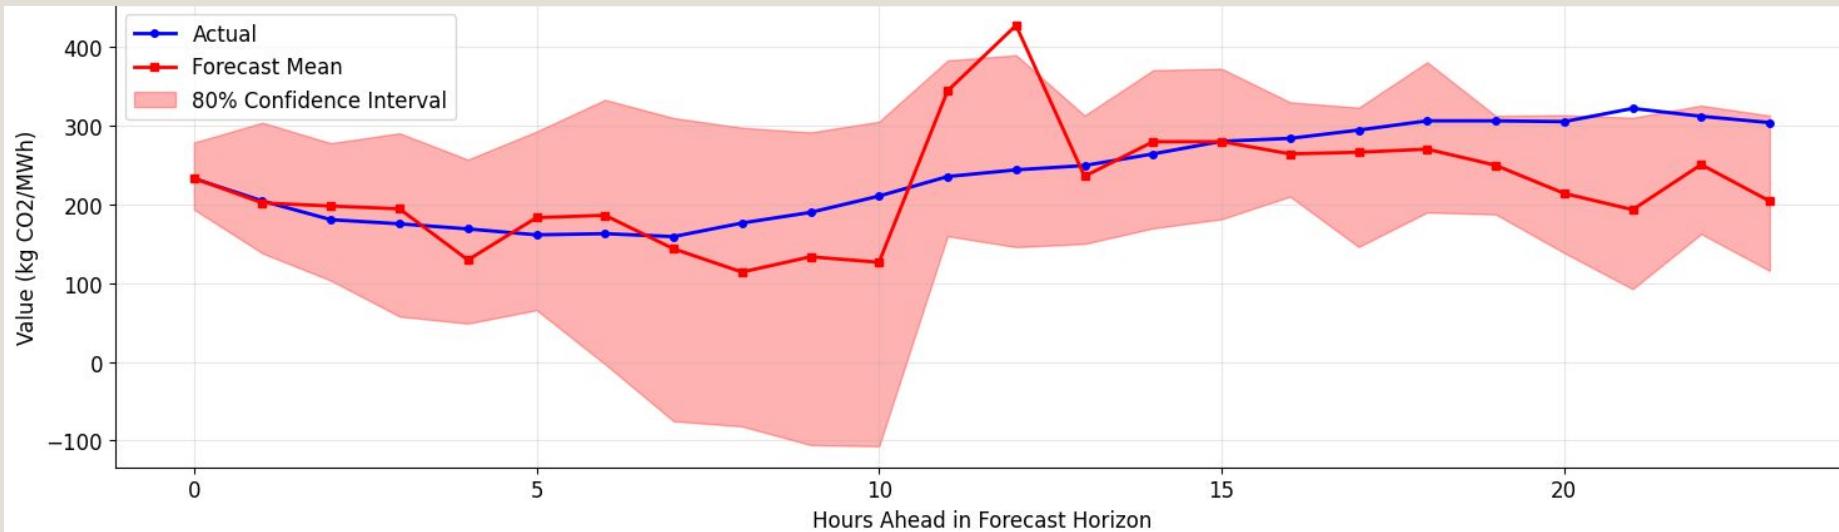
- Trained on 27 time series diverse datasets comprising energy, transportation and economics.
- Training corpus comprised of 7965 univariate time series totalling 352 million tokens - Open Source.

Why Foundation Model?

- Represent a new frontier, derived from the success of LLM.
- Powerful generalization capabilities across various tasks and domains, often with minimal or no fine-tuning.
- Achieve scalability and adaptability in time series analysis.

Zero Shot Forecasting

Forecast for Carbon Intensity (kg CO₂/MWh)

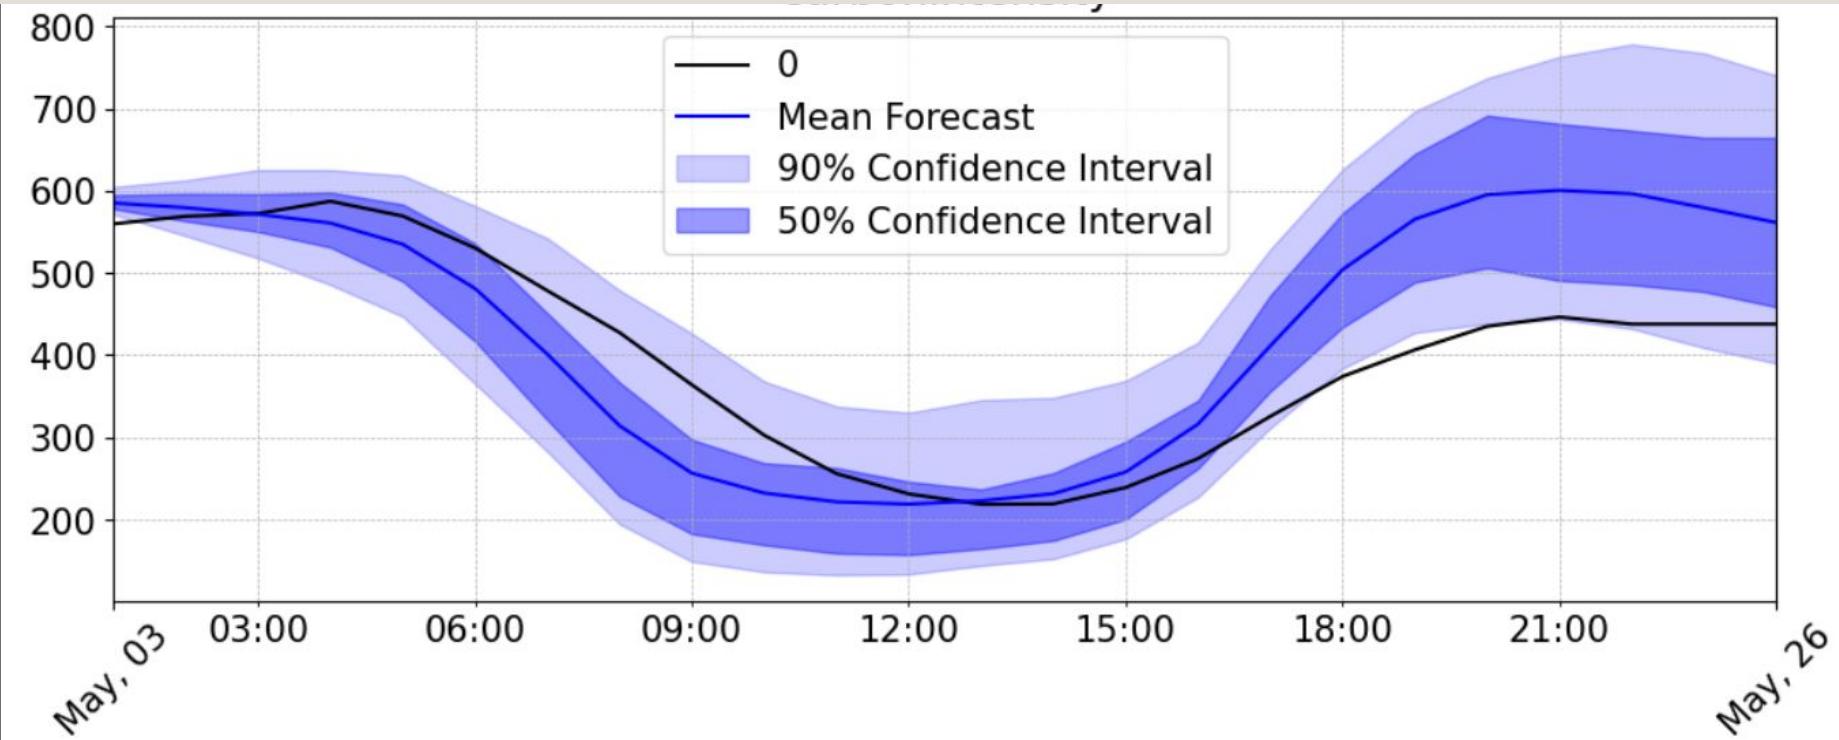


MAE: 49.364

RMSE: 66.988

MAPE: 20.4%

Preliminary Fine Tuning Forecasting



RMSE: 109.520

MAPE: 60.4%

Observations

Zero Shot:

- Model has no idea about the data.
- The model successfully captured the daily pattern (the "shape" of the curve) right away, showing it understands the fundamental drivers of carbon intensity.

Fine Tuning:

- Model have been trained for 20 Epochs on the data.
- Performed worse than zero-shot.
- Direction is correct but very high magnitude error.
- Too early for any meaningful prediction.

Conclusion & Future Remark - Lag Llama

Conclusion:

- **The Foundation Model is Powerful:** Lag-Llama's "zero-shot" performance (20% error) was surprisingly accurate and set a very strong baseline. Its prior experience is highly relevant to our data.
- **Initial Fine-Tuning Did Not Beat the Baseline:** Our first attempt at fine-tuning resulted in a higher average error than the zero-shot test. This can happen if the training is too short or if the model struggles to learn multiple, complex patterns at once.
- **Pattern vs. Magnitude:** In both tests, the model did a great job of predicting the *pattern* (the daily ups and downs). The main challenge is improving its accuracy in predicting the *magnitude* (the actual values).

Future Work:

- The task could be a challenge for future groups.

Outlook

- Finalize report
- Make small adjustments to the models, e.g. tweak the selection of the generated artificial features
- Further improving the loss function
- Trying to understand better the behavior of the model at various inputs

Thank You

References

- [1] K. Rasul et al. (2023), *Lag-Llama: Towards Foundation Models for Probabilistic Time Series Forecasting*, arXiv, [arXiv:2310.08278](https://arxiv.org/abs/2310.08278).